

Car_Price_Prediction

```
In [47]: #importing Libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [48]: df = pd.read_csv('CarPrice.csv') #readig dataset
```

```
In [49]: df.head()
```

Out[49]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	fuelsystem	borerat
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.4
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.4
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.6
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.1
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.1

5 rows × 26 columns

```
In [50]: df.tail()
```

Out[50]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	fuelsystem	boreratio
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78

5 rows × 26 columns

```
In [51]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling             205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize             205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke               205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm              205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg           205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [52]: df.describe()
```

```
Out[52]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horse
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.0
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.329756	3.255415	10.142537	104.1
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.270844	0.313597	3.972040	39.5
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.540000	2.070000	7.000000	48.0
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.150000	3.110000	8.600000	70.0
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.310000	3.290000	9.000000	95.0
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.580000	3.410000	9.400000	116.0
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	3.940000	4.170000	23.000000	288.0

```
In [54]: df.isnull().sum()
```

```
Out[54]:
```

car_ID	0
symboling	0
CarName	0
fueltype	0
aspiration	0
doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	0
carlength	0
carwidth	0
carheight	0
curbweight	0
enginetype	0
cylindernumber	0
enginesize	0
fuelsystem	0
boreratio	0
stroke	0
compressionratio	0
horsepower	0
peakrpm	0
citympg	0
highwaympg	0
price	0
dtype: int64	

```
In [55]: df.duplicated().sum()
```

```
Out[55]: 0
```

```
In [56]: df.shape #give number of rows and number of columns
```

```
Out[56]: (205, 26)
```

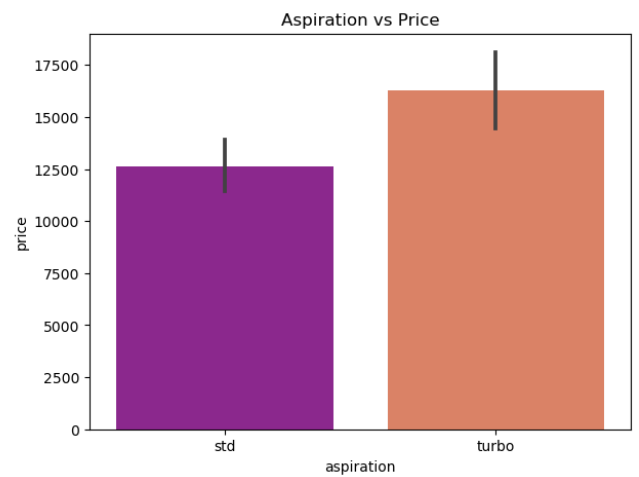
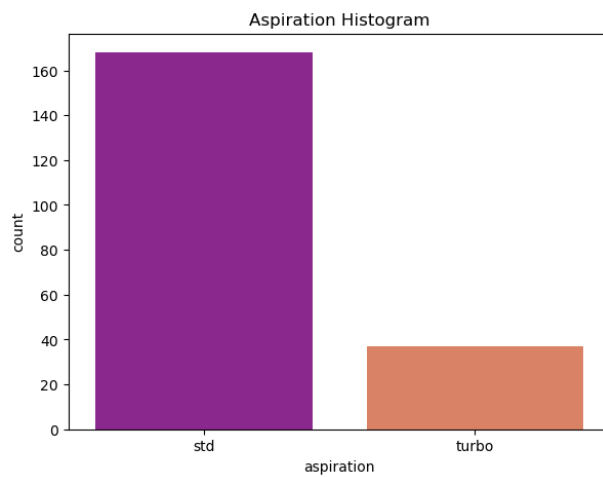

```
In [57]: print(df.price.describe(percentiles=[0.225,0.50,0.75,0.85,0.98,1]))
```

```
count      205.000000
mean     13276.710571
std       7988.852332
min       5118.000000
22.5%     7609.000000
50%      10295.000000
75%      16503.000000
85%      18500.000000
98%      36809.600000
100%     45400.000000
max       45400.000000
Name: price, dtype: float64
```

Data Visualization

```
In [59]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.title("Door number Histogram")
sns.countplot(data=df, x='doornumber', palette="plasma")
plt.subplot(1,2,2)
plt.title('Door number vs Price')
sns.barplot(data=df, x='doornumber', y='price', palette="plasma")
plt.show()

plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.title("Aspiration Histogram")
sns.countplot(data=df, x='aspiration', palette="plasma")
plt.subplot(1,2,2)
plt.title('Aspiration vs Price')
sns.barplot(data=df, x='aspiration', y='price', palette="plasma")
plt.show()
```



```

In [31]: #Fuel type
colors=sns.color_palette('pastel')
labels=df['fueltype'].dropna().unique()
plt.figure(figsize=(18,10))
plt.subplot(1,2,1)


plt.title('fueltype_percentage')
plt.pie(df['fueltype'].value_counts(),labels=labels,colors=colors,autopct='%0.2f%%')
plt.subplot(1,2,2)
plt.title('fueltype bar chart')
sns.countplot(x='fueltype',data=df,palette=colors)
df.fueltype.value_counts(dropna=False)

```



```

Out[31]: gas      185
         diesel    20
         Name: fueltype, dtype: int64

```



```
In [32]: dff=pd.DataFrame(df.groupby(['fueltype'])['price'].mean().sort_values(ascending=False))
dff.plot.bar()
plt.title("Fuel type vs Average price")
plt.show()
dff=pd.DataFrame(df.groupby(['carbody'])['price'].mean().sort_values(ascending=False))
dff.plot.bar()
plt.title("Car type vs Average price")
plt.show()
```



```
In [33]: predict="price"
df=df[['symboling','wheelbase','carwidth', 'carheight', 'curbweight', 'enginesize','boreratio', 'stroke','compressionratio',
```

```
In [34]: x=np.array(df.drop([predict],1))

y=np.array(df[predict])
```

C:\Users\baps\AppData\Local\Temp\ipykernel_13096\3277193174.py:1: FutureWarning: In a future version of pandas all argument
s of DataFrame.drop except for the argument 'labels' will be keyword-only.
x=np.array(df.drop([predict],1))

```
In [35]: print(x)
print(y)

[[ 3.000e+00  8.860e+01  6.410e+01 ... 5.000e+03  2.100e+01  2.700e+01]
 [ 3.000e+00  8.860e+01  6.410e+01 ... 5.000e+03  2.100e+01  2.700e+01]
 [ 1.000e+00  9.450e+01  6.550e+01 ... 5.000e+03  1.900e+01  2.600e+01]
 ...
 [-1.000e+00  1.091e+02  6.890e+01 ... 5.500e+03  1.800e+01  2.300e+01]
 [-1.000e+00  1.091e+02  6.890e+01 ... 4.800e+03  2.600e+01  2.700e+01]
 [-1.000e+00  1.091e+02  6.890e+01 ... 5.400e+03  1.900e+01  2.500e+01]]

[[13495.  16500.  16500.  13950.  17450.  15250.  17710.
 18920.  23875.  17859.167 16430.  16925.  20970.  21105.
 24565.  30760.  41315.  36880.  5151.  6295.  6575.
 5572.  6377.  7957.  6229.  6692.  7609.  8558.
 8921.  12964.  6479.  6855.  5399.  6529.  7129.
 7295.  7295.  7895.  9095.  8845.  10295.  12945.
 10345.  6785.  8916.5  8916.5  11048.  32250.  35550.
 36000.  5195.  6095.  6795.  6695.  7395.  10945.
 11845.  13645.  15645.  8845.  8495.  10595.  10245.
 10795.  11245.  18280.  18344.  25552.  28248.  28176.
 31600.  34184.  35056.  40960.  45400.  16503.  5389.
 6189.  6669.  7689.  9959.  8499.  12629.  14869.
 14489.  6989.  8189.  9279.  9279.  5499.  7099.
 6649.  6849.  7349.  7299.  7799.  7499.  7999.
 8249.  8949.  9549.  13499.  14399.  13499.  17199.
 19699.  18399.  11900.  13200.  12440.  13860.  15580.
 16900.  16695.  17075.  16630.  17950.  18150.  5572.
 7957.  6229.  6692.  7609.  8921.  12764.  22018.
 32528.  34028.  37028.  31400.5  9295.  9895.  11850.
 12170.  15040.  15510.  18150.  18620.  5118.  7053.
 7603.  7126.  7775.  9960.  9233.  11259.  7463.
 10198.  8013.  11694.  5348.  6338.  6488.  6918.
 7898.  8778.  6938.  7198.  7898.  7788.  7738.
 8358.  9258.  8058.  8238.  9298.  9538.  8449.
 9639.  9989.  11199.  11549.  17669.  8948.  10698.
 9988.  10898.  11248.  16558.  15998.  15690.  15750.
 7775.  7975.  7995.  8195.  8495.  9495.  9995.
 11595.  9980.  13295.  13845.  12290.  12940.  13415.
 15985.  16515.  18420.  18950.  16845.  19045.  21485.
 22470.  22625. ]]
```

```
In [36]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=100)
```

Random Forest model

```
In [37]: from sklearn.ensemble import RandomForestRegressor
```

```
In [38]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=100)
print('training data shape is:{}'.format(x_train.shape))
print('training label shape is:{}'.format(y_train.shape))
print('testing data shape is:{}'.format(x_test.shape))
print('testing data shape is:{}'.format(y_test.shape))

training data shape is:(164, 13).
training label shape is:(164,).
testing data shape is:(41, 13).
testing data shape is:(41,).
```

```
In [39]: from sklearn.ensemble import RandomForestRegressor
regressor=RandomForestRegressor()
```

```
In [40]: regressor.fit(x,y)
```

```
Out[40]: RandomForestRegressor
RandomForestRegressor()
```

```
In [41]: regressor.score(x_train,y_train)
```

```
Out[41]: 0.9889896872489429
```

```
In [42]: regressor.score(x_test,y_test)
```

```
Out[42]: 0.9881855875703067
```

```
In [43]: from sklearn.metrics import accuracy_score
predictions=regressor.predict(x_test)
```

```
In [44]: percentage=regressor.score(x_test,y_test)
percentage
```

```
Out[44]: 0.9881855875703067
```

```
In [60]: print(regressor.score(x_train,y_train))  
print(f"test set:{len(x_test)}")  
print(f"Accuracy={percentage*100}%")
```

0.9889896872489429

test set:41

Accuracy=98.81855875703067%

THANK YOU!

GitHub: <https://github.com/anujtiwari21?tab=repositories> (<https://github.com/anujtiwari21?tab=repositories>)